

GETTING STARTED WITH NoSQL

Aaron Benton / @bentonam



- The problem with SQL
- History of NoSQL
- Database theories
- Modeling
- Patterns

THE PROBLEM WITH SQL

- Designed to run on large servers
- Built for Vertical Scaling
- Separated Models i.e. Tables

IN DEVELOPMENT

- We assemble objects as a whole
 - Cart
 - Order
 - Profile
 - Product
- Saving objects requires
 - Deconstructing
 - Multiple rows
 - Multiple tables

IMPEDANCE MISMATCH

"The object-relational impedance mismatch is a set of conceptual and technical difficulties that are often encountered when a relational database management system (RDBMS) is being used by a program written in an object-oriented programming language or style, particularly when objects or class definitions are mapped in a straightforward way to database tables or relational schema." - Wikipedia

**YOU CAN IGNORE
SQL**



**AND TRUST THE ORM
MAGIC**

memegenerator.net



Google BigTable

Amazon DynamoDB

NoSQL

No SQL

Not only SQL



#nosql

WHAT IS NOSQL?

- non-relational
- cluster friendly
- generally open-source
- 21st century
- schema-less

TYPES OF NOSQL DATABASES

- Key-Value: Redis, Riak, Memcached
- Column-family: Cassandra, HBase, BigTable
- Document: CouchDB, Couchbase, MongoDB
- Graph: Neo4J, Giraph, OrientD

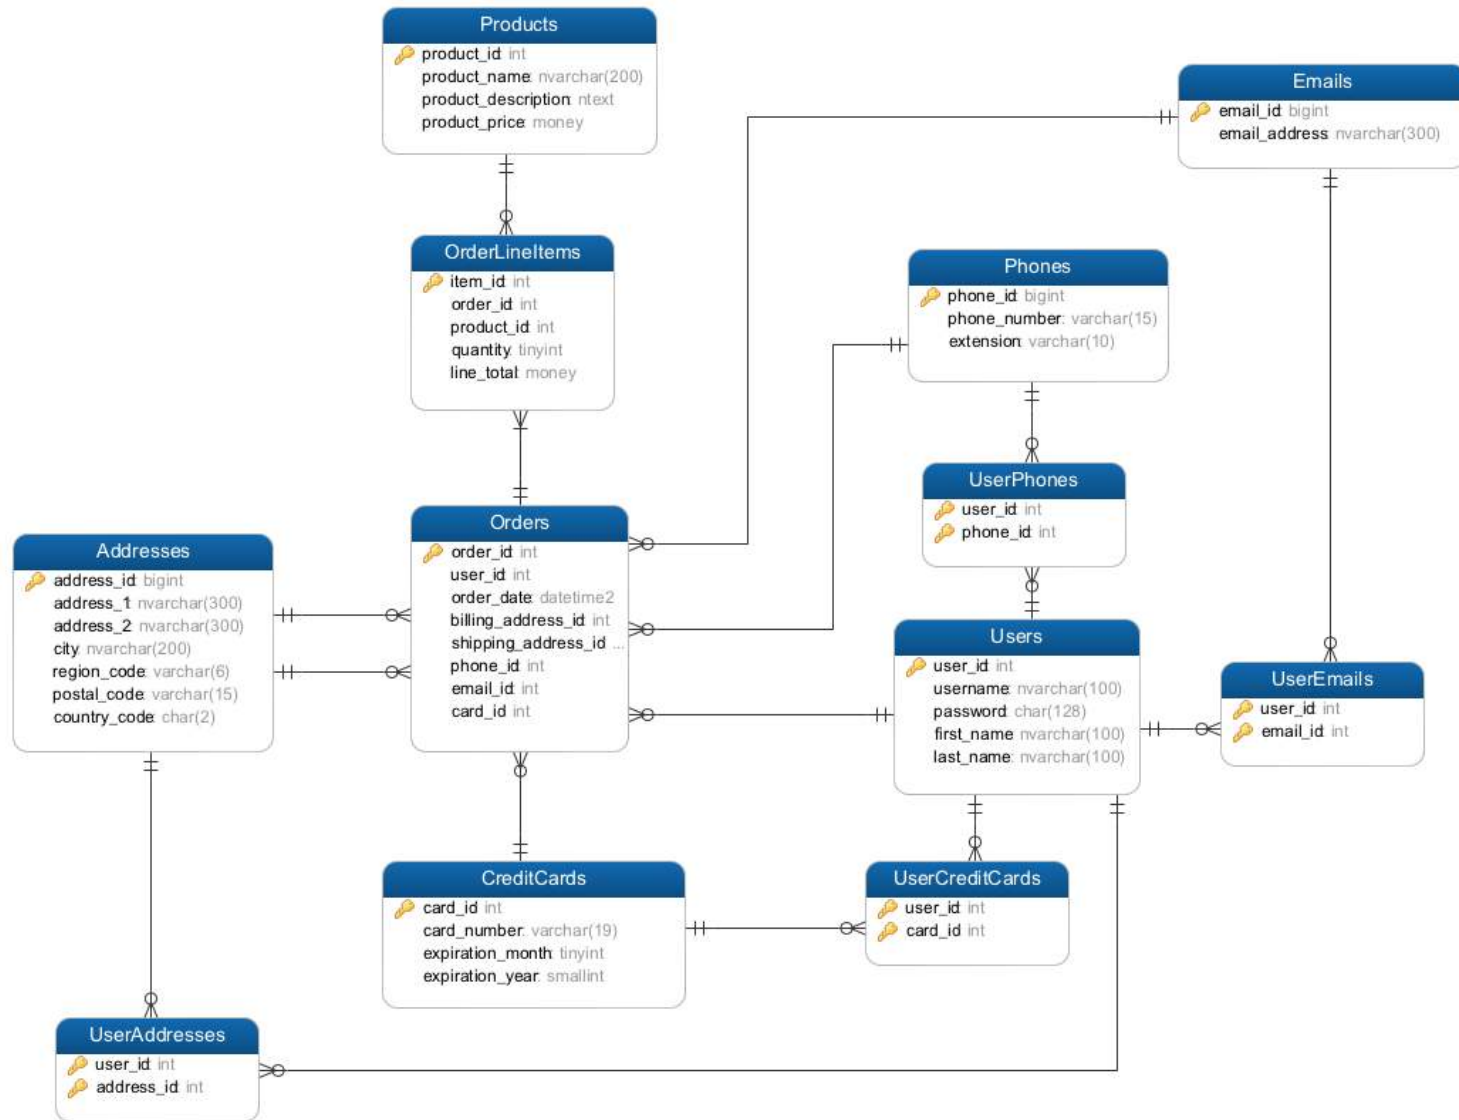
AGGREGATES

```
// Order.cfc
component accessors="true"{
    property="order_id" type="numeric";
    property="order_date" type="date";
    property="products" type="array";
    property="user_id" type="numeric";
    property="billing_address_1" type="string";
    property="billing_address_2" type="string";
    property="billing_city" type="string";
    property="billing_region_code" type="string";
    property="billing_postal_code" type="string";
    property="billing_country_code" type="string";
    property="shipping_address_1" type="string";
    property="shipping_address_2" type="string";
    property="shipping_city" type="string";
    property="shipping_region_code" type="string";
    property="shipping_postal_code" type="string";
    property="shipping_country_code" type="string";
    property="card_number" type="string";
    property="expiration_month" type="numeric";
    property="expiration_year" type="numeric";
}
```

```
// Order.cfc
component accessors="true"{
    property="order_id" type="numeric";
    property="order_date" type="date";
    property="products" type="array";
    property="user" type="User";
    property="billing" type="Address";
    property="shipping" type="Address";
    property="cc_info" type="CreditCard";
}
```

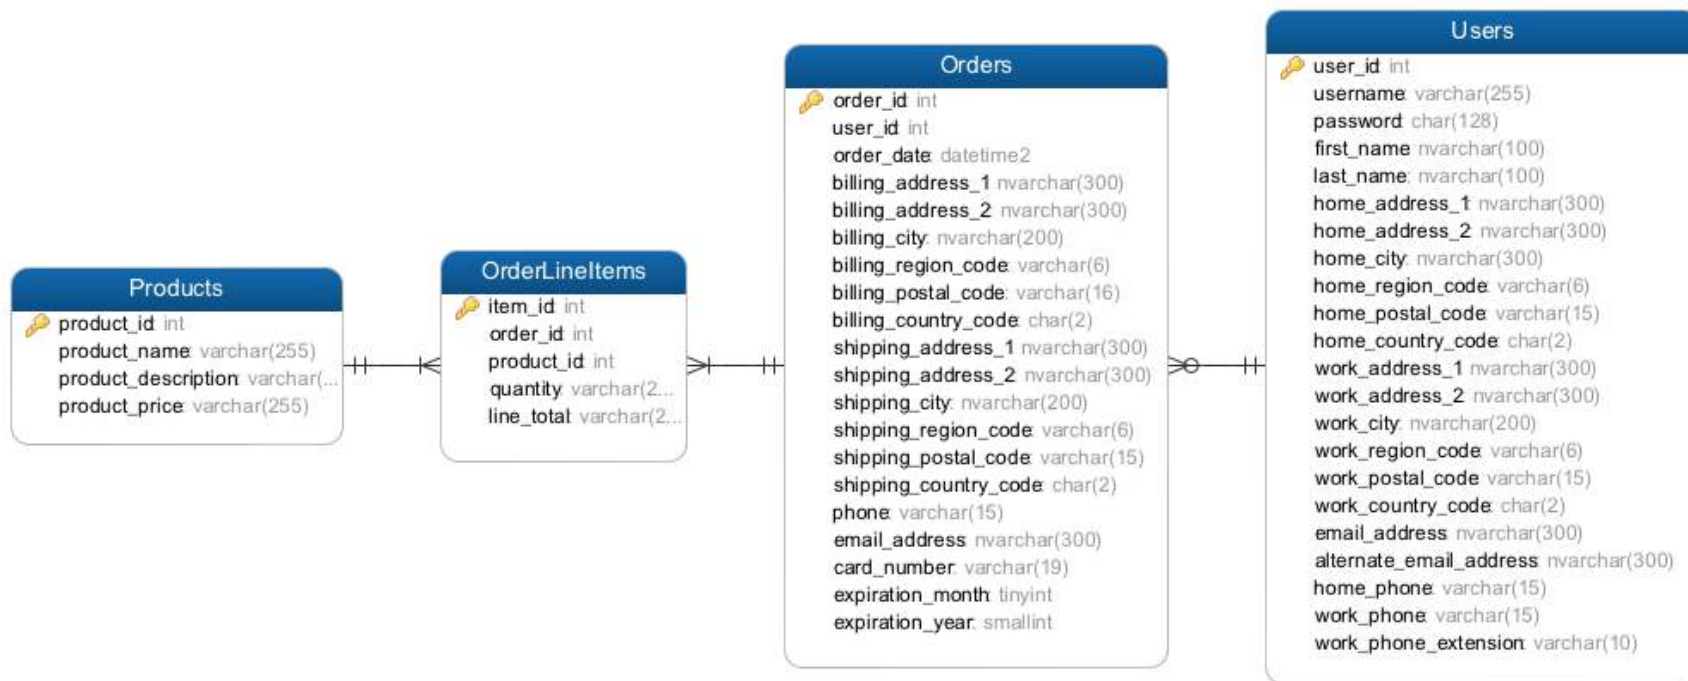
NORMALIZATION

- Minimize data redundancy
- Structured models
- Logical queries
- Fast inserts / updates
- Less storage requirements

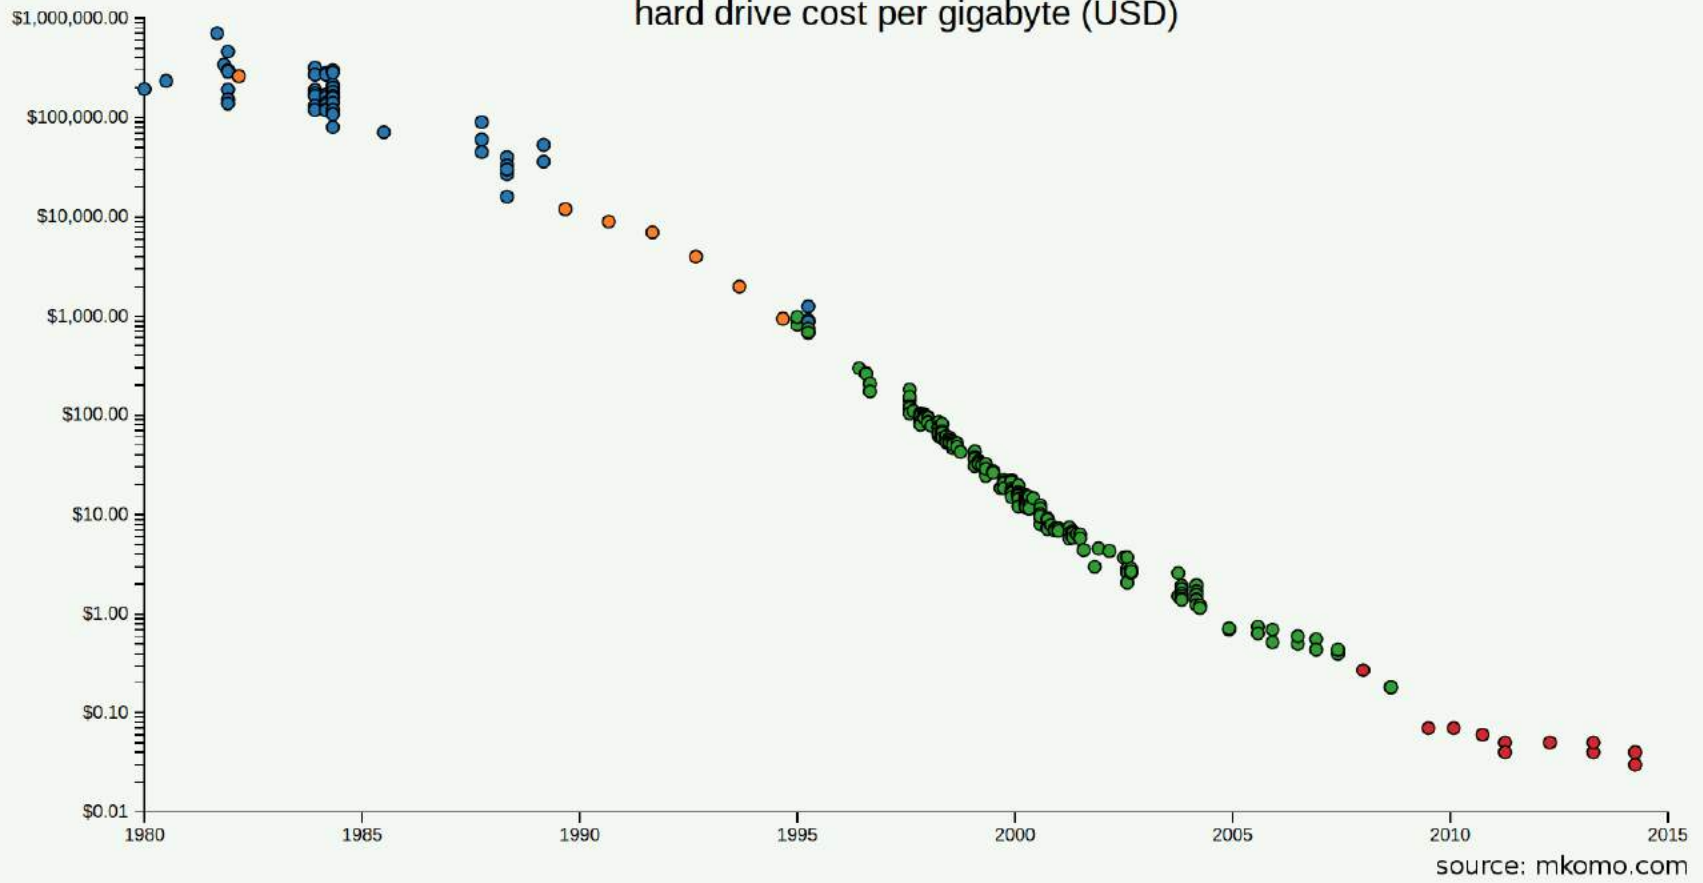


DENORMALIZATION

- Minimize JOINS
- Fast reads
- Repeated data
- More storage



hard drive cost per gigabyte (USD)



TRANSACTION PROCESSING

ACID

- Atomicity
- Consistency
- Isolation
- Durability

BASE

- Basically Available
- Soft State
- Eventual Consistency

ATOMIC TRANSACTIONS

Aggregate Orientated == NoSQL - Graph

CONFLICT RESOLUTION

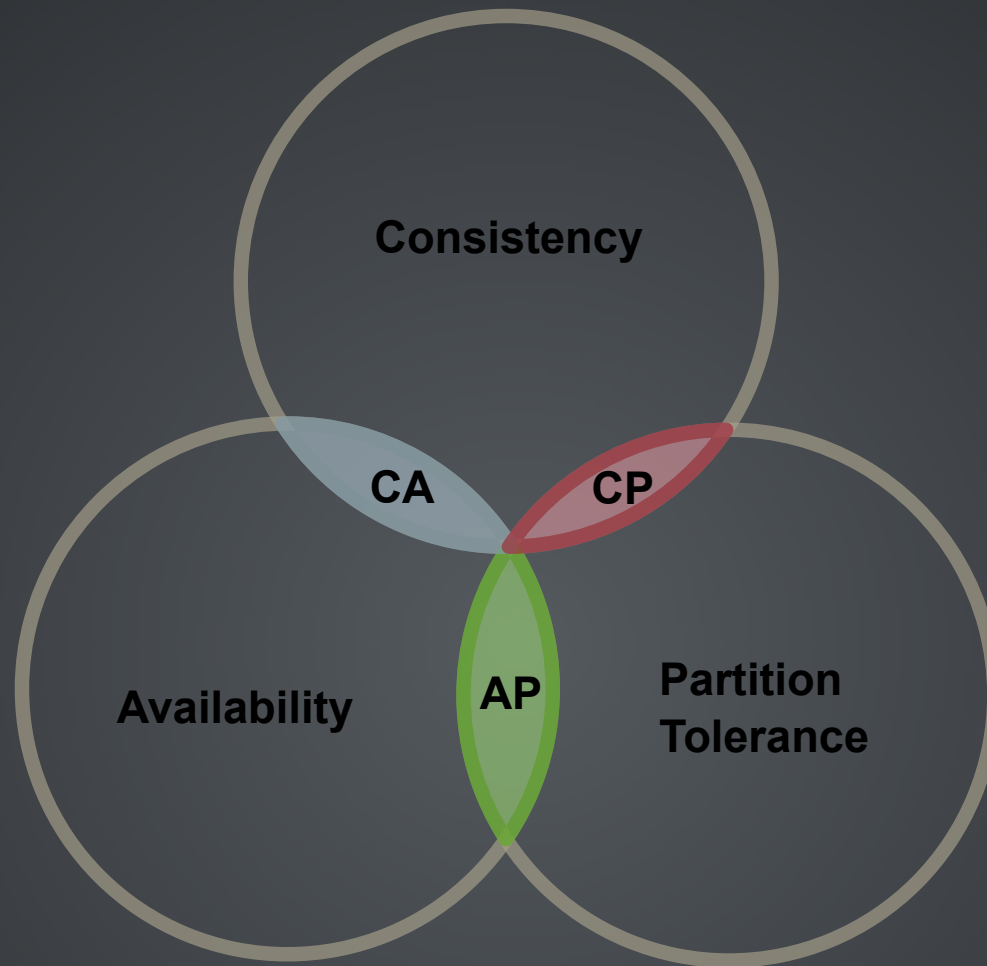
CONSISTENCY

SHARDING VS. REPLICATION

CAP THEOREM

- Consistency
- Availability
- Partition Tolerance

You can only provide 2 of the 3



RDBMS SCHEMAS

- Known Models
- Fixed Fields
- Data types
- Database managed
- Change can be difficult



NOSQL SCHEMAS

- Any type of data
- Flexible
- Application managed
- Change is easy



IMPLICIT SCHEMA

~~SCHEMA LESS~~

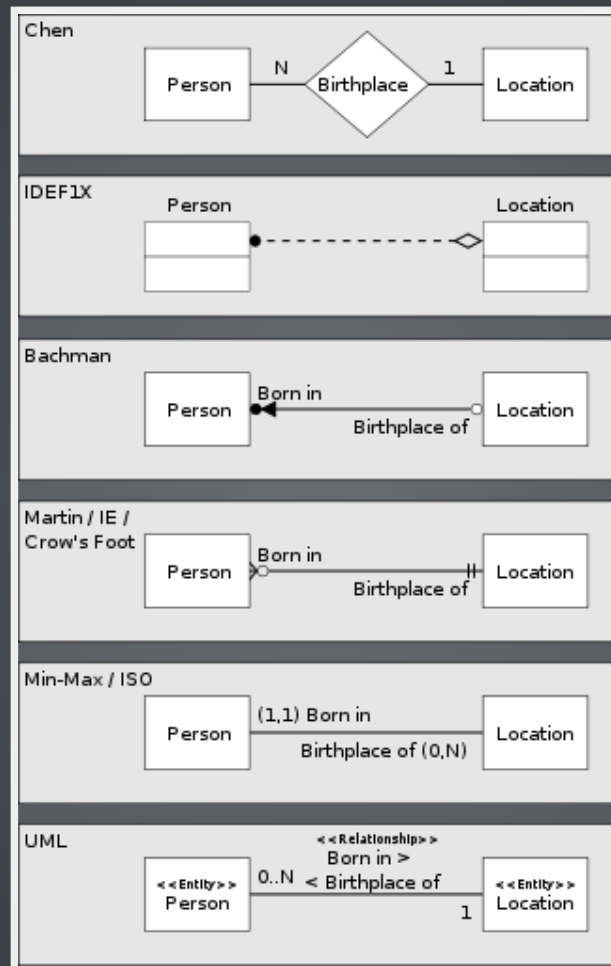
DATA / ENTITY RELATIONSHIP MODELING

- Conceptual Data Model
- Logical Data Model
- Physical Data Model

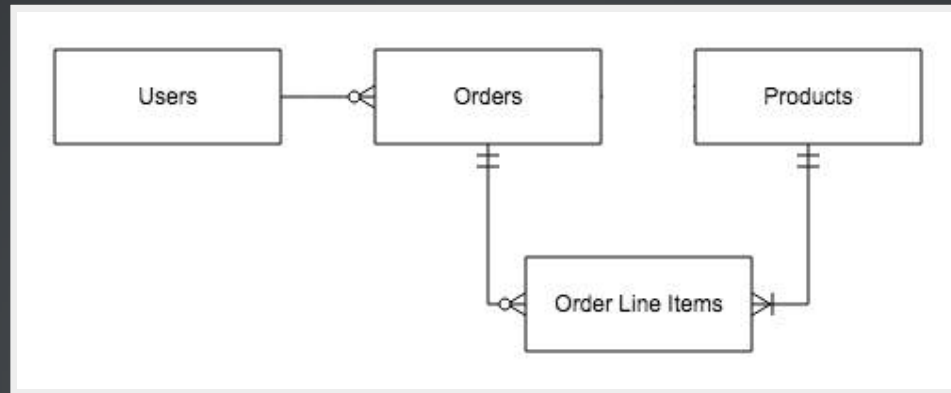
CONCEPTUAL DATA MODEL

- Entity Names
- Entity Relationships

MODELING NOTATIONS



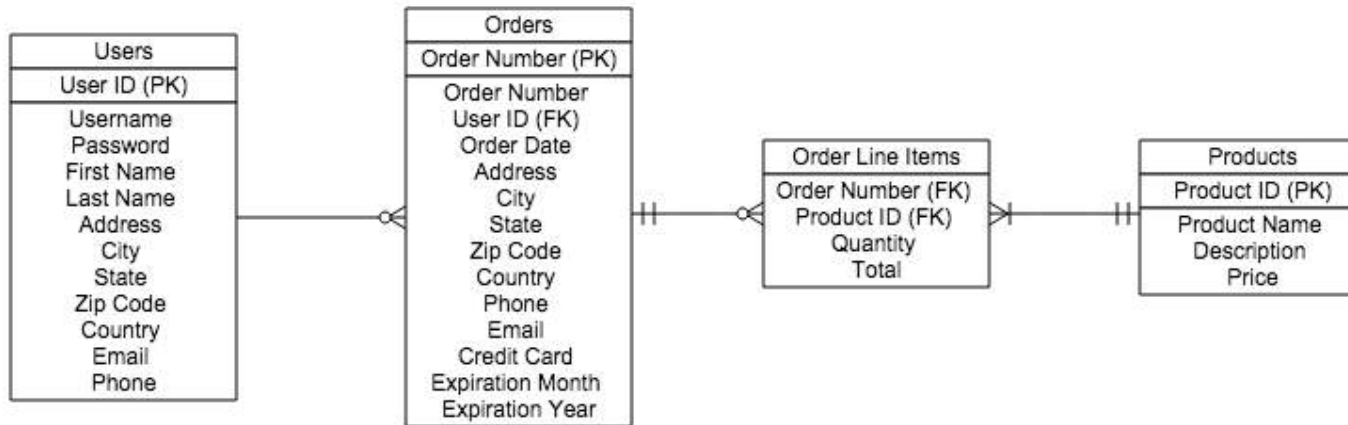
CONCEPTUAL DATA MODEL



LOGICAL DATA MODEL

- Entity Names
- Entity Relationships
- Attributes
- Primary / Foreign Keys

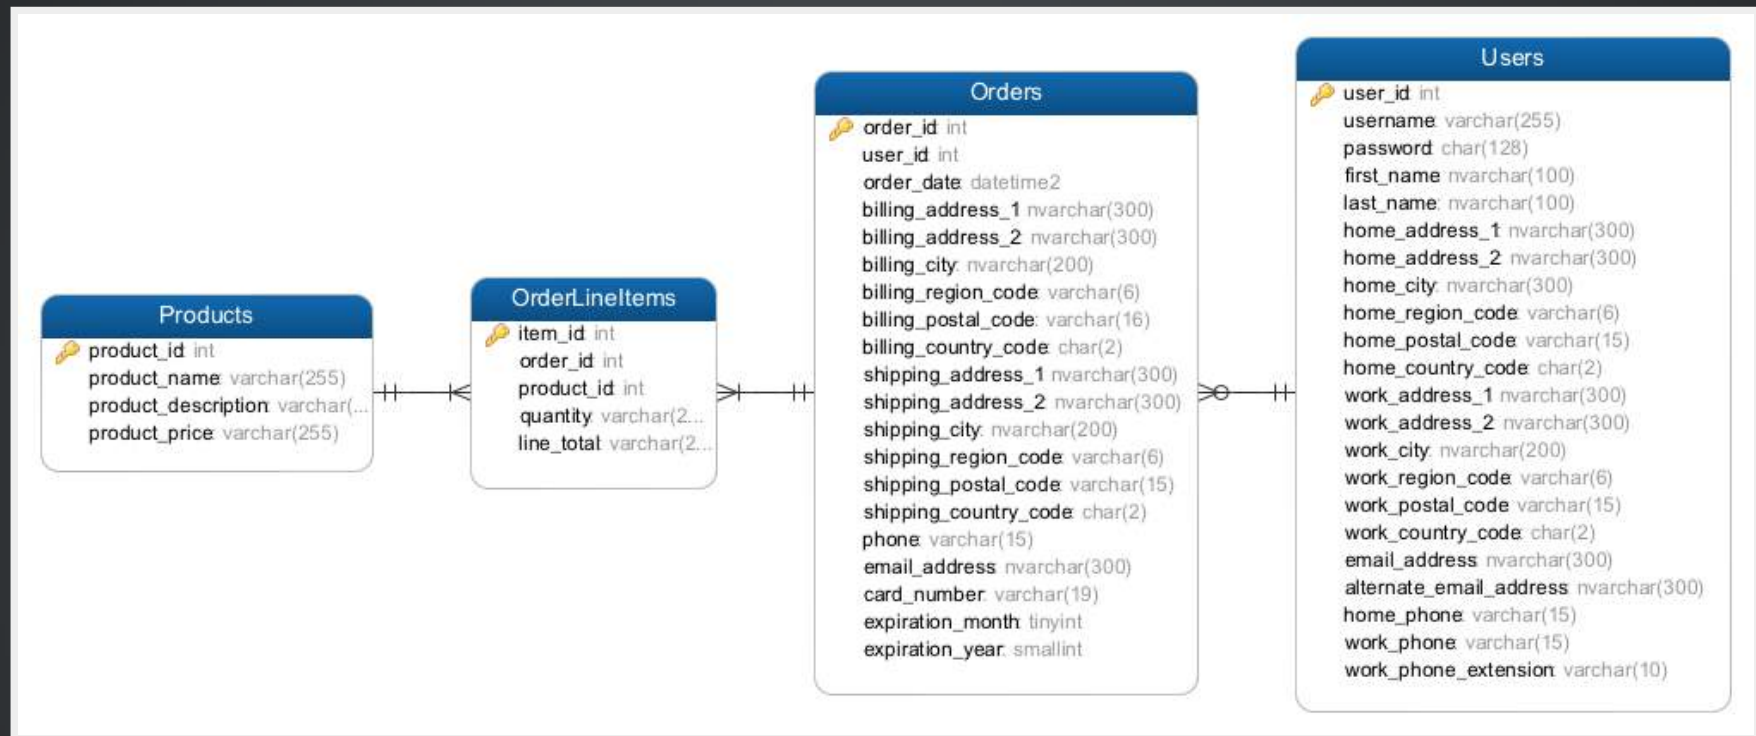
LOGICAL DATA MODEL



PHYSICAL DATA MODEL

- Entity -> Table Names
- Attributes -> Field Names
- Keys -> Primary / Foreign Keys
- Data Types

PHYSICAL DATA MODEL



PHYSICAL DATA MODEL IN NOSQL



```
{  
  "user_id": 123,  
  "username": "jdoe",  
  "first_name": "John",  
  "last_name": "Doe",  
  "email": "john.doe@mail.com",  
  "password": "88142f883cba2b527fdbbc60a943b899"  
}
```



```
type: object
properties:
  _id:
    type: string
    description: The ID of the document
  user_id:
    type: integer
    description: > An auto-incremented number from `users_counter`
    that is the ID of the user
  first_name:
    type: string
    description: The users first name
  last_name:
    type: string
    description: The users last name
  username:
    type: string
    description: A unique username chosen by the user
  friends:
    type: array
    description: An array of user_id who the user is friends with
  created_on:
    type: integer
    description: An epoch time in seconds when the user was created
```

KEY DESIGN

- Prefixing
- Predictable
- Counter ID
- Unpredictable
- Combinations

PREFIXING

- user_123
- u::john.doe@mail.com
- user-123
- user_123_orders
- order_123
- o::john.doe@mail.com
- product-123
- user_123_orders

PREDICTABLE

Key: user_john.doe@mail.com

Key: user_jdoe

```
JSON
{
  user_id : 123
  username : "jdoe"
  first_name : "John"
  last_name : "Doe"
  email : "john.doe@mail.com"
  password : "88142f883cba2b527fdbbc60a943b899"
```

```
JSON
{
  user_id : 123
  username : "jdoe"
  first_name : "John"
  last_name : "Doe"
  email : "john.doe@mail.com"
  password : "88142f883cba2b527fdbbc60a943b899"
```

COUNTER ID

Key: user_123

```
JSON
{
  user_id : 123
  username : "jdoe"
  first_name : "John"
  last_name : "Doe"
  email : "john.doe@mail.com"
  password : "88142f883cba2b527fdbbc60a943b899"
```

Key: user_counter

```
JSON
{
  counter : 414
```

UNPREDICTABLE

Key: 23ad6bac-7599-4874-af98-7af734027834



COMBINATIONS

- user_123_preferences
- user_jdoe_order_23ad6bac-7599-4874-af98-7af734027834
- user_john.doe@mail.com_comment_5664

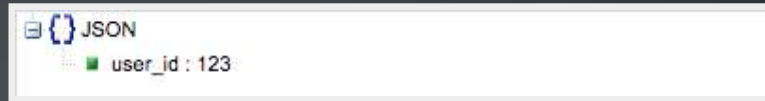
DOCUMENT PATTERNS

LOOKUP PATTERN

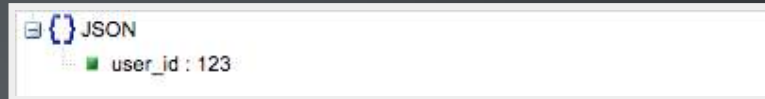
Key: user_123



Key: user_john.doe@mail.com



Key: user_jdoe

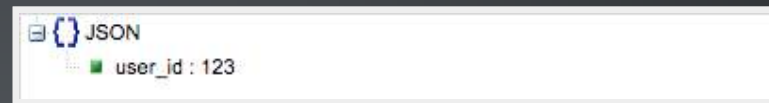


LOOKUP PATTERN FOR AUTH

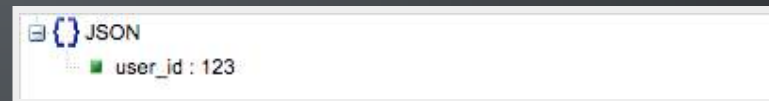
Key: user_123



Key: user_john.doe@mail.com_
88142f883cba2b527fdbbc60a943b8



Key: user_jdoe_
88142f883cba2b527fdbbc60a943b8



EMBEDDING

```
... ■ username : "jdoe"
... ■ first_name : "John"
... ■ last_name : "Doe"
... ■ email : "john.doe@mail.com"
... ■ password : "88142f883cba2b527fdbbc60a943b899"
... [] addresses
...   [] 0
...     ■ address_1 : "123 Missing St."
...     ■ address_2 : ""
...     ■ city : "Greensboro"
...     ■ region_code : "NC"
...     ■ postal_code : "27409"
...     ■ type : "home"
...   [] 1
...     ■ address_1 : "321 Smithfield Ave."
...     ■ address_2 : ""
...     ■ city : "Greensboro"
...     ■ region_code : "NC"
...     ■ postal_code : "27409"
...     ■ type : "work"
```

REFERRING / LINKING

Key: user_123

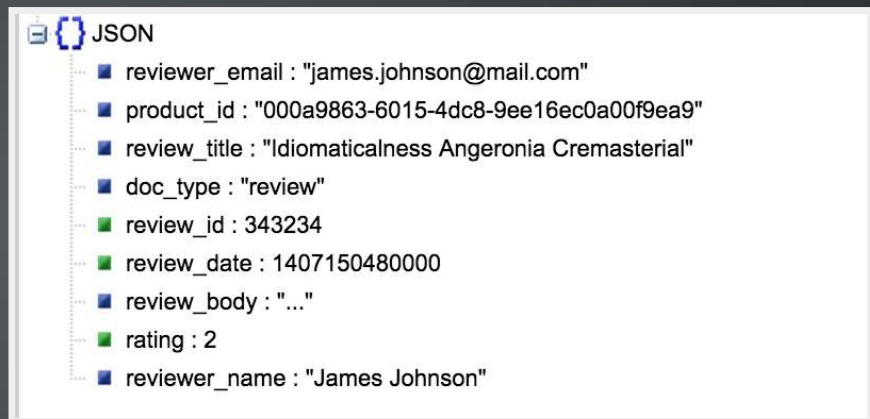
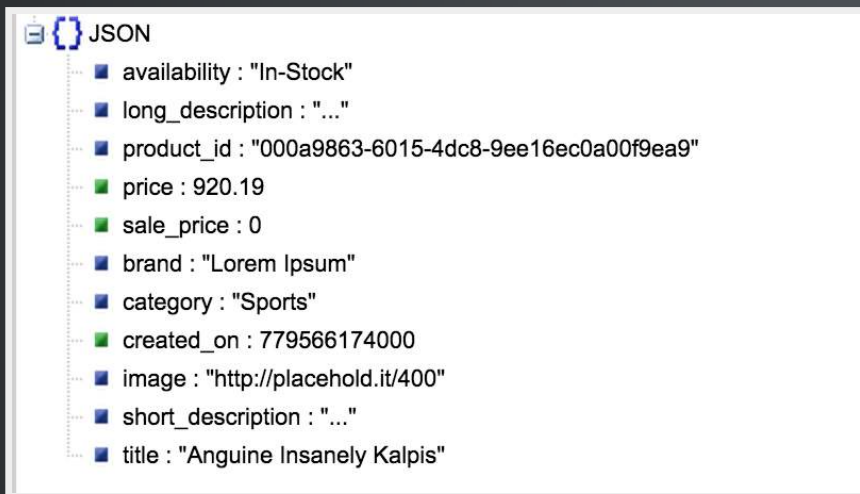
```
JSON
{
  user_id : 123
  username : "jdoe"
  first_name : "John"
  last_name : "Doe"
  email : "john.doe@mail.com"
  password : "88142f883cba2b527fdbbc60a943b899"
  friends : [
    0 : 456
    1 : 743
    2 : 834
  ]
}
```

Key: user_456

```
JSON
{
  user_id : 456
  username : "jsmith"
  first_name : "Jane"
  last_name : "Smith"
  email : "jane.smith@mail.com"
  password : "87a48ef776efbbe92651257bc1a52e84"
  friends : [
    0 : 123
    1 : 654
    2 : 837
  ]
}
```

PARENT-REFERENCING

Key: product_000a9863-6015-4dc8-9ee16ec0a00f9ea9 Key: product_000a9863-6015-4dc8-9ee16ec0a00f9ea9_review_343234



FAKEIT

Inspired by [Swagger](#), fakeit is a CLI data generator based on YAML models that outputs JSON, YAML, CSON, or CSV formats

```
npm install fakeit -g
```

- <https://www.npmjs.com/package/fakeit>
- <https://github.com/bentonam/fakeit>
- <https://github.com/bentonam/fakeit-examples>

FAKEIT

- Generate fixed or random number of documents per model
- Event Transforms: Pre / Post Run, Pre / Post Build
- Data generation via FakerJS, ChanceJS, Custom or Static
- Data Typing
- Model Dependencies
- JSON, YAML, CSON, CSV output formats
- Output to File, Zip, Couchbase or Sync Gateway

```
name: Users
type: object
key: _id
data:
  min: 200
  max: 500
  pre_run: >
    globals.user_counter = 0;
properties:
  id:
    type: string
    data:
      post_build: "return 'user_' + this.user_id;"
  type:
    type: string
    data:
      value: "user"
  user_id:
    type: integer
    data:
      build: "return ++globals.user_counter;"
  name:
    type: string
    data:
      fake: "{{name.firstName}} {{name.lastName}}"
  phone:
    type: string
    data:
      build: "return chance.phone();"
  created_on:
    type: string
    data:
      fake: "{{date.past}}"
      post_build: "return new Date(this.created_on).toISOString();"
```


FAKEIT DEMO

Is SQL Going Away?

NO

CONSIDERATIONS

- How do you work with your data?
- Do you work with the same aggregates all the time?
- What are you trying to achieve?
- Where are you starting at?
- Do you need finite data and highly complex relationships?
- Is the tabular structure working for you?
- Do you want to scale vertically or horizontally?
- Does your data need to be data centralized or decentralized?

QUESTIONS?

SLIDES AVAILABLE AT:

bit.ly/gsw-nosql

RESOURCES

- [Introduction to NoSQL by Martin Fowler](#)
- [Relationships are Hard NoSQL Data Modeling by Curt Gratz](#)
- [Workshop: NoSQL Data Modelling by Jan Steemann](#)
- [CAP Twelve Years Later: How the "Rules" Have Changed by Eric Brewer](#)
- [NoSQL Databases: An Overview by Pramod Sadalage](#)